SINCE AI HACKATHON 2026 · FONECTA CHALLENGE

Connect Suggests – Customer Approves

No customer walks to a closed door. And no owner has to log in to stop it.

Flovr · Alex, Vova and Zahar (Flovr's AI agent) · synthetic Fonecta dataset only

THE PROBLEM

Owners don't log in. So the hours go wrong exactly when it matters.

53%

of the 2025 pilot suggestions got **no reply at all** (1,610 of 3,031).

Reply rate by channel, same pilot:

- **65%** WhatsApp
- **50%** SMS
- **40%** Email
- **32%** App push

~40%

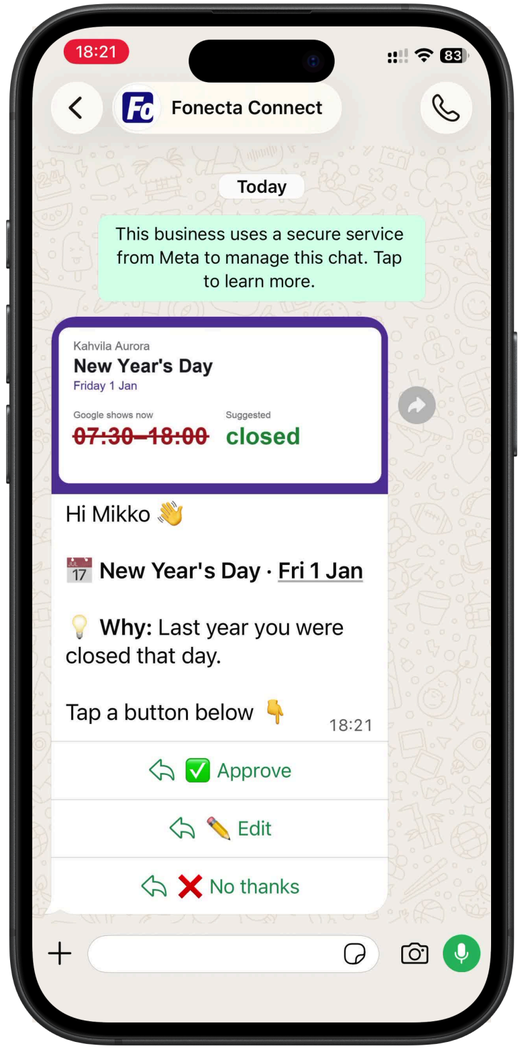
of needed changes have **no usable history** to copy (Fonecta's own note). That's where it gets hard.

Source: pilot_decision_log_2025.csv (3,031 suggestions) and the starter kit README, in the challenge dataset.

THE SOLUTION

The whole chain, working

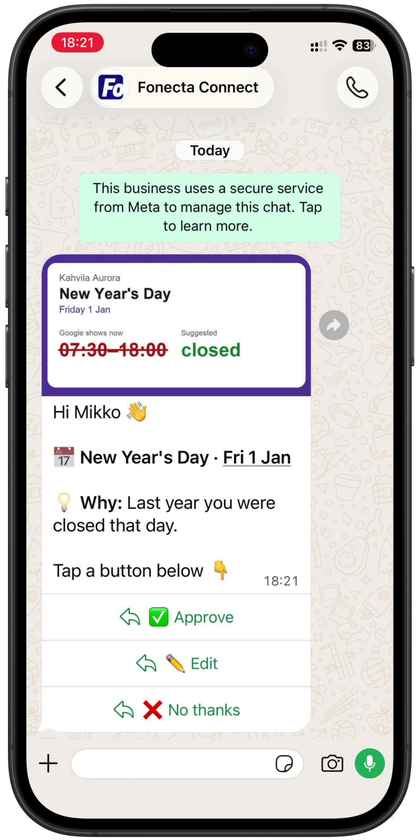
1. **Detect:** a model over every business × special day: own history, shopping centre, peers, last year's activity, closed-door reviews, Google user edits.
2. **Suggest:** the hours candidate with the best track record. Confident → suggest, unsure → ask, weak → stay quiet.
3. **Contact:** on WhatsApp, in the owner's language (fi/sv/nb/en). One message per date: what Google shows now, what we suggest, and why.
4. **Decide:** no login, no app. Approve / Edit / No thanks buttons. Edits are typed or spoken in the chat ("9 to 13").
5. **Apply:** written to Yext, then a confirmation of where it shows and how soon.



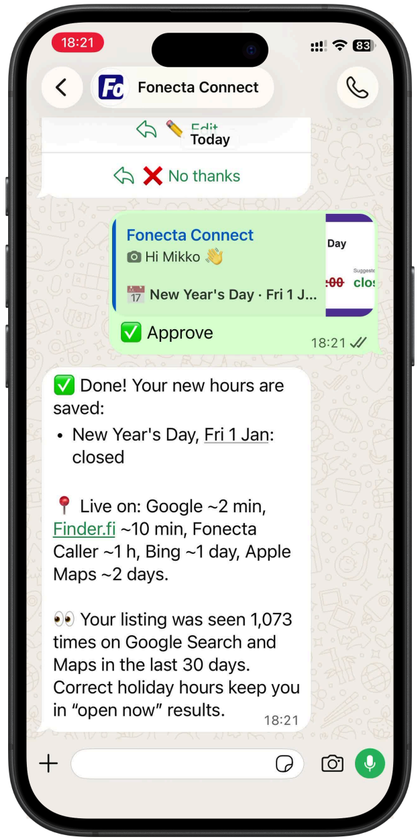
Real iPhone screen: our WhatsApp test number, synthetic business from the Fonecta dataset.

THE CUSTOMER EXPERIENCE

The owner never leaves WhatsApp



Ask: one date, the reason, three buttons



Confirm: live on Google in ~2 min, and their real search views



Call: no answer in 30 min → a WhatsApp call

Real iPhone screen recordings from the running app, 7 Nov. Free text uses Claude Haiku (Gemini fallback); voice notes are transcribed locally with Whisper.

RELIABILITY

The customer always has the last word

- Nothing reaches Yext without an owner decision, a standing permission (public holidays, confident only, followed by an inform message with one-tap **Undo**) or an undo. That's enforced in code and unit-tested.
- Unsure → we **ask**, never assume (new customer with no history: only confident days are suggested, the rest become questions).
- Signs the business has **shut down for good** (recent reviews saying "permanently closed", or Google users marking it closed) → we don't send holiday-hours messages to a business that may not exist; instead it's **flagged for Connect staff** to check the listing.
- Hours already right → **silence**.
- No answer in **30 min** → Connect **calls the owner** on WhatsApp (the owner can change the 30 min in the chat), and tries once more 30 min later. **Only if there's still no answer** → the next approver gets the same message and a call. Whoever taps first decides; a call never approves anything.
- Every suggestion keeps its reasons, the message and the answer: a full **audit trail**.

HOW IT LEARNS

Every answer makes the next suggestion better

- **Edit** → remembered for that day, and the pattern (e.g. +1 h) applies to similar days.
- **"Open as usual" / "Not this year"** → that day goes quiet.
- **No reason given** → lower confidence: next time we ask instead of suggest.
- **Standing permission** → applied, then the owner is informed.
- 2025 pilot decisions seed each owner's history from day one.

What Connect has learned about this owner

Kahvila Aurora · FI-D0001 · Connect staff console

- No 2025 pilot history: new to suggestions, so we ask rather than assume.
- Standing permission for public holidays: no.
- Declined INDEPENDENCE_DAY (Not this year, 2026-11-07) → quiet for the 2026 season.
- Edited XMAS_EVE: 07:30-13:00 → 08:00-14:00 (2026-11-07) → next time we suggest 08:00-14:00.
- Edit pattern: opening +30 min, closing +60 min vs our suggestions (1 edit(s)) → applied to similar suggestions.
- Declined BOXING_DAY without a clear reason → confidence -0.2, we ask instead of suggesting.
- Approved 1 suggestion(s) as sent → slightly more confident next time.

Real output of the app's learning module (app/learning.py), as listed in the Connect staff console, after 4 demo answers: edit Christmas Eve, "Not this year" for Independence Day, approve New Year's Day, decline Boxing Day with no reason. Owners never see this page.

RESULTS · FONECTA'S SCORING SCRIPT

Better than the baseline, on businesses the model never saw

	Score	Exact precision	Recall	Harmful	2025 back-test
Baseline (kit)	54.5	0.90	0.59	224	53.2
Ours, first model	61.8	0.84	0.70	396	59.3
Ours, final model	**65.5**	0.81	**0.75**	404	**63.5**

The final model finds far more of the wrong hours (recall 0.75 vs the baseline's 0.59), at the cost of more wrong-state guesses in the raw scored file (404 vs 224). In the product, anything under 0.6 confidence becomes a **question** rather than a suggestion, and nothing is published without the owner.

Fonecta's own scoring script (evaluate.py) on the 803 businesses they gave answers for (dev 2026), each scored by a model trained without it; thresholds picked on the other season. 2025 back-test = the same on last season. The official score is Fonecta's, on their hidden businesses.

EXTENSIBILITY + LIVE DATA

One loop, many jobs

- **Reviews:** a bad review gets a reply drafted in its language → owner posts / edits / skips → Yext Reviews API. Plus a weekly summary of **every** review, good, mixed and bad: the count, the average, the split, what people praised **and** what they complained about.
- **Live events** (bonus): Helsinki LinkedEvents, live. A concert 650 m away ends 20:15 → "Stay open till 21:15?". 5 of 106 Helsinki restaurants got one this weekend. Cafés get "open earlier?" for morning events, and anyone gets a busy-day heads-up for staff and stock. No attendance data is the gap.
- **Channel-agnostic:** one decision engine, any channel. Live today: WhatsApp (buttons, chat edits, voice notes) and WhatsApp voice calls. Built and switched off: SMS (for owners without smartphones) and iMessage.
- **Decision log:** every suggestion, answer and change in one table, per business: who, what, when, which channel.
- **Staff shifts** (synthetic schedule) and **watched pages** ("watch this page for events near my café"): the same ask-then-apply loop.

PRIVACY, SECURITY AND AI

Responsible by default

- **Data:** only the synthetic Fonecta dataset and public Helsinki events. No real customer data.
- **Secrets:** Yext, WhatsApp and Gemini keys live in the keychain and env vars. None in the repo, its history, the video or this deck.
- **Control:** no write without the owner (enforced in code, unit-tested). Buttons carry signed, expiring tokens; owners never open a web page.
- **Messaging:** WhatsApp only after the owner writes first or with an approved template; calls only with call permission; opt-outs respected.
- **AI in the product:** Claude Haiku reads free-text replies (Gemini fallback). Its output is validated, and it never decides. Whisper runs locally, so audio stays on our machine.
- **Suggestions** come from a classic model with reasons we can show, not a chatbot.
- **Built with AI:** Zahar (Flovr's agent, Claude Code) wrote most of the code; every result was checked by running it.
- **Pre-existing:** built on Fonecta's starter kit (MIT); everything else was built at the event. Zahar's own platform stays out of the repo.

THANK YOU

Connect Suggests – Customer Approves

Next: a WhatsApp-first pilot before Christmas, measured against 2025's 47% reply rate.

About 600 Finnish Connect customers · shadow mode 23 Nov · live from 1 Dec · Sweden and Norway in 2027 (pilot plan attached)

github.com/zahar-io/fonecta (public) · live demo: demo.zahar.io · demo video: 2 min, real app

Flovr · Alex, Vova and Zahar (Flovr's AI agent) · flovr.co.uk